Logic For First Submission

Overall Architecture:

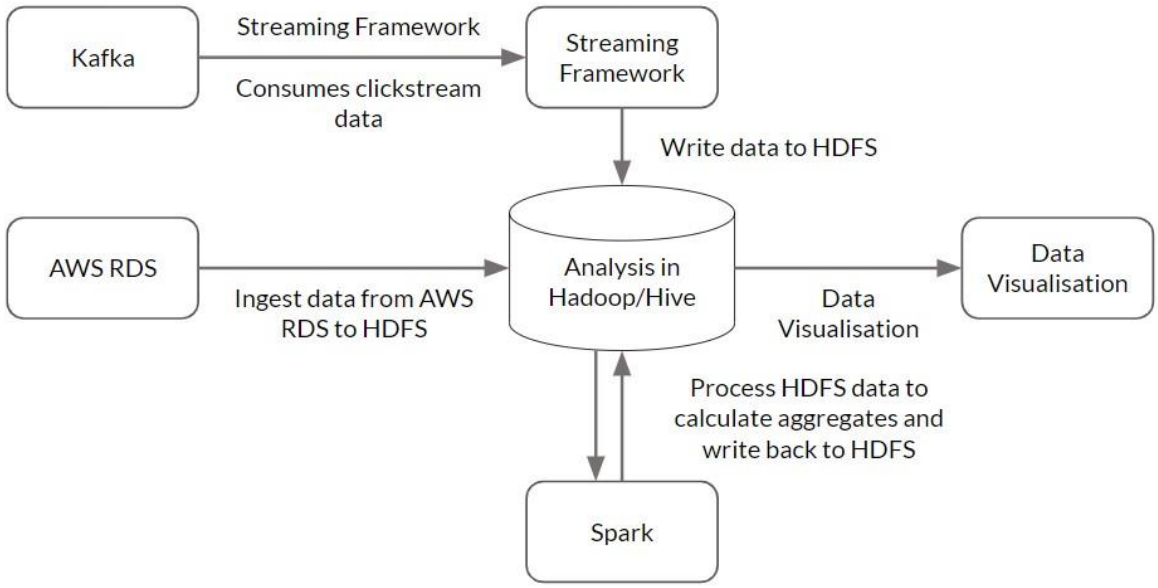

According to the given architecture, we have performed the following steps:

In the first step, we ingest stream data from Kafka into a stream processing framework. Then, we consume batch data from AWS RDS into HDFS. Where data aggregation is necessary, we read and write our data from Spark to HDFS. The final aggregated data is stored in Hive tables. The data in Hive tables is later queried to calculate different metrics.

For the first submission (MID), we consume data from Kafka and RDS and load them into Hive managed tables. This process is split into four tasks:

Task 1: Write a job to consume clickstream data from Kafka and ingest to Hadoop

Please find commented file "**spark_kafka_to_local.py**"

1) At first, we create a Spark session with application name "Kafka to HDFS"

```
spark = SparkSession.builder \
    .appName("Kafka To HDFS") \
    .getOrCreate()
spark.sparkContext.setLogLevel('WARN')
```

2) Then we read the incoming clickstream data from Kafka and storing the data in the hdfs. To do this, we use the credentials provided on the platform.

Details of the Kafka Broker are as follows:

- **Bootstrap-server** - 18.211.252.152
- **Port** - 9092
- **Topic** - de-capstone3

The code to read the data is shown below:

```
df = spark \
    .readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "18.211.252.152:9092") \
    .option("subscribe","de-capstone3") \
    .option("startingOffsets","earliest") \
    .load()
df.printSchema()
```

3) Then we transform the data by dropping a few columns and changing **value** column data type

```
df = df.withColumn('value_str', df['value'].cast('string').alias('key_str')).drop('value') \
       .drop('key','topic','partition','offset','timestamp','timestampType')
```

4) We write the data to hdfs directory and keep it running until terminated (Write stream in json format).

```
df1=df.writeStream \
  .outputMode("append") \
  .format("json") \
  .option("truncate","false") \
  .option("path", "/home/hadoop/clickstream_data") \
  .option("checkpointLocation","/home/hadoop/clickstream_checkpoint") \
  .start()

console_df = df \
    .writeStream \
    .format("console") \
    .outputMode("append") \
```

Code execution output:
After running the code, we get the following output

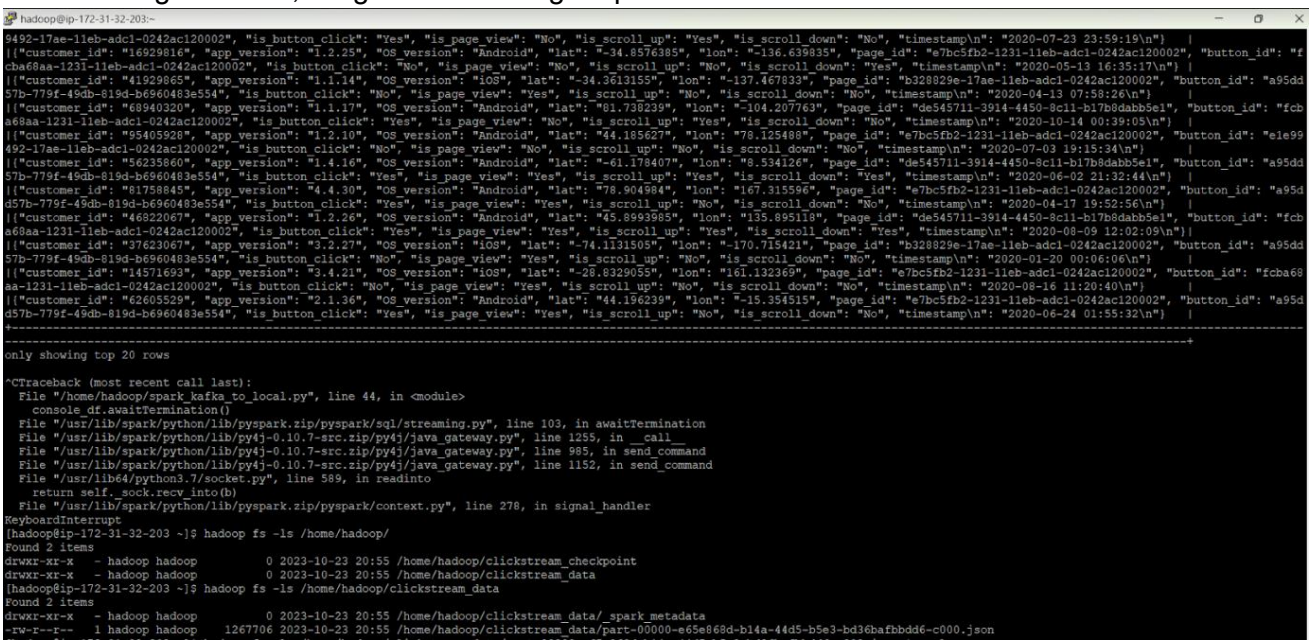

We now clean the loaded Kafka data to a more structured format and save it to a csv file.

Please find commented file "**spark_local_flatten.py**"

1) Firstly, we get the existing spark session "Kafka to HDFS"

```
spark = SparkSession.builder \
    .appName("Kafka To HDFS") \
    .getOrCreate()
spark.sparkContext.setLogLevel('WARN')
```

2) Then we read the clickstream data that we stored in json format into the dataframe

```
df = spark.read.json('/home/hadoop/clickstream_data/part*')
```

3) We extract the columns from json **value_str** in dataframe and create a new dataframe with new columns

```
df1 = df.select(
    get_json_object(df["value_str"],"$.customer_id").alias("customer_id"),
    get_json_object(df["value_str"],"$.app_version").alias("app_version"),
    get_json_object(df["value_str"],"$.OS_version").alias("OS_version"),
    get_json_object(df["value_str"],"$.lat").alias("lat"),
    get_json_object(df["value_str"],"$.lon").alias("lon"),
    get_json_object(df["value_str"],"$.page_id").alias("page_id"),
    get_json_object(df["value_str"],"$.button_id").alias("button_id"),
    get_json_object(df["value_str"],"$.is_button_click").alias("is_button_click"),
    get_json_object(df["value_str"],"$.is_page_view").alias("is_page_view"),
    get_json_object(df["value_str"],"$.is_scroll_up").alias("is_scroll_up"),
    get_json_object(df["value_str"],"$.is_scroll_down").alias("is_scroll_down"),
    get_json_object(df["value_str"],"$.timestamp\n").alias("click_timestamp")
)
```

4) We then save the new dataframe to a csv file in hdfs directory

```
df1.coalesce(1).write.format('com.databricks.spark.csv').mode('overwrite').save('/home/hadoop/clickstream_data_flatten', header = 'false')
```

The steps to run the python files and the steps is provided in commented file "**step1.pdf**"

Code execution output:
After running the code, we get the following output

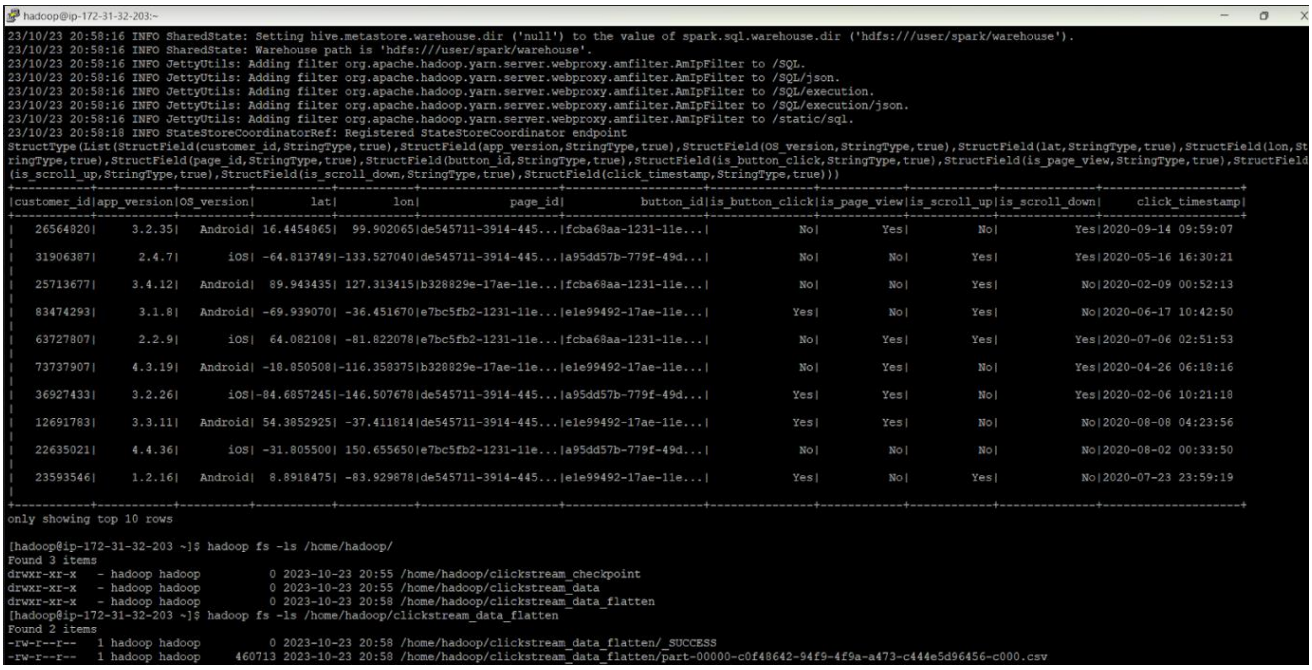

```
hadoop@ip-172-31-32-203:~
23/10/23 20:58:16 INFO SharedState: Setting hive.metastore.warehouse.dir ('null') to the value of spark.sql.warehouse.dir ('hdfs:///user/spark/warehouse').
23/10/23 20:58:16 INFO SharedState: Warehouse path is 'hdfs:///user/spark/warehouse'.
23/10/23 20:58:16 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL.
23/10/23 20:58:16 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL/json.
23/10/23 20:58:16 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL/execution.
23/10/23 20:58:16 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL/execution/json.
23/10/23 20:58:16 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /static/sql.
23/10/23 20:58:18 INFO StateStoreCoordinatorRef: Registered StateStoreCoordinator endpoint
StructType(List(StructField(customer_id,StringType,true),StructField(app_version,StringType,true),StructField(OS_version,StringType,true),StructField(lat,StringType,true),StructField(lon,St
ringType,true),StructField(page_id,StringType,true),StructField(button_id,StringType,true),StructField(is_button_click,StringType,true),StructField(is_page_view,StringType,true),StructField
(is_scroll_up,StringType,true),StructField(is_scroll_down,StringType,true),StructField(click_timestamp,StringType,true)))
+-----------+-----------+----------+-----------+-----------+--------------------+--------------------+--------------+------------+------------+--------------+--------------------+
|customer_id|app_version|OS_version|        lat|        lon|             page_id|           button_id|is_button_click|is_page_view|is_scroll_up|is_scroll_down|     click_timestamp|
+-----------+-----------+----------+-----------+-----------+--------------------+--------------------+--------------+------------+------------+--------------+--------------------+
|   26564820|      3.2.35|   Android| 16.4454865|  99.902065|de545711-3914-445...|fcba68aa-1231-11e...|            No|         Yes|          No|           Yes|2020-09-14 09:59:07|
|   31906387|      2.4.7|       iOS| -64.813749|-133.527040|de545711-3914-445...|a95dd57b-779f-49d...|            No|          No|         Yes|           Yes|2020-05-16 16:30:21|
|   25713677|      3.4.12|   Android|  89.943435| 127.313415|b328829e-17ae-11e...|fcba68aa-1231-11e...|            No|          No|         Yes|            No|2020-02-09 00:52:13|
|   83474293|      3.1.8|   Android| -69.939070| -36.451670|e7bc5fb2-1231-11e...|e1e99492-17ae-11e...|           Yes|          No|         Yes|            No|2020-06-17 10:42:50|
|   63727807|      2.2.9|       iOS|  64.082108| -81.822078|e7bc5fb2-1231-11e...|fcba68aa-1231-11e...|            No|         Yes|         Yes|           Yes|2020-07-06 02:51:53|
|   73737907|      4.3.19|   Android| -10.850508|-116.358375|b328829e-17ae-11e...|e1e99492-17ae-11e...|            No|         Yes|          No|           Yes|2020-04-26 06:18:16|
|   36927433|      3.2.26|       iOS|-84.6857245|-146.507678|de545711-3914-445...|a95dd57b-779f-49d...|           Yes|         Yes|          No|           Yes|2020-02-06 10:21:18|
|   12691783|      3.3.11|   Android| 54.3852925| -37.411814|de545711-3914-445...|e1e99492-17ae-11e...|           Yes|         Yes|          No|            No|2020-08-08 04:23:56|
|   22635021|      4.4.36|       iOS| -31.805500| 150.655650|e7bc5fb2-1231-11e...|a95dd57b-779f-49d...|            No|          No|          No|            No|2020-08-02 00:33:50|
|   23593546|      1.2.16|   Android|  8.8918475| -83.929878|de545711-3914-445...|e1e99492-17ae-11e...|           Yes|          No|         Yes|            No|2020-07-23 23:59:19|
+-----------+-----------+----------+-----------+-----------+--------------------+--------------------+--------------+------------+------------+--------------+--------------------+
only showing top 10 rows

[hadoop@ip-172-31-32-203 ~]$ hadoop fs -ls /home/hadoop/
Found 3 items
drwxr-xr-x   - hadoop hadoop          0 2023-10-23 20:55 /home/hadoop/clickstream_checkpoint
drwxr-xr-x   - hadoop hadoop          0 2023-10-23 20:55 /home/hadoop/clickstream_data
drwxr-xr-x   - hadoop hadoop          0 2023-10-23 20:58 /home/hadoop/clickstream_data_flatten
[hadoop@ip-172-31-32-203 ~]$ hadoop fs -ls /home/hadoop/clickstream_data_flatten
Found 2 items
-rw-r--r--   1 hadoop hadoop          0 2023-10-23 20:58 /home/hadoop/clickstream_data_flatten/_SUCCESS
-rw-r--r--   1 hadoop hadoop     460713 2023-10-23 20:58 /home/hadoop/clickstream_data_flatten/part-00000-c0f48642-94f9-4f9a-a473-c444e5d96456-c000.csv
```

Task 2: Write a job to ingest the bookings data from AWS RDS to Hadoop

Please find commented file "**step2.pdf**" for the sqoop command

- To do this, we write a sqoop job using the credentials provided on the platform

The RDS connection string and credentials are as follows:

- RDS **Connection String** -

  `jdbc:mysql://upgraddetest.cyaielc9bmnf.us-east-1.rds.amazonaws.com/testdatabase`

- **Username** - student
- **Password** - STUDENT123
- **Table Name** - bookings

```
sqoop import \
--connect jdbc:mysql://upgraddetest.cyaielc9bmnf.us-east-1.rds.amazonaws.com/testdatabase \
--table bookings \
--username student --password STUDENT123 \
--target-dir /home/hadoop/bookings_data \
-m1
```

Code execution output:
After running the code, we get the following output

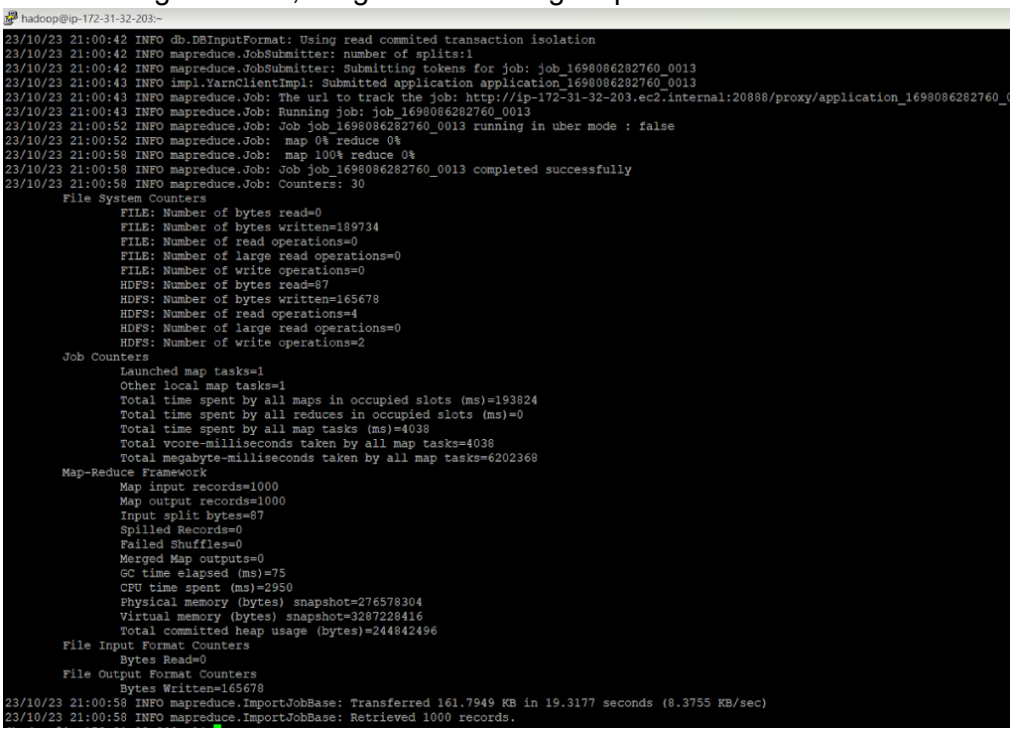

<u>Task 3</u>: Create aggregates for finding total bookings date-wise using Spark script

To prepare aggregates, data is read from HDFS, processed by a processing framework such as Spark and written back to HDFS to create a Hive table for the aggregated data.

Please find commented file "**datewise_bookings_aggregates_spark.py**"

1) To perform aggregation of data, we first create a Spark session with application name "datewise_booking"

```
spark = SparkSession.builder.appName("datewise_booking").getOrCreate()
sc = spark.sparkContext
sc
```

2) Then we read the bookings data into the dataframe

```
df = spark.read.csv("/home/hadoop/bookings_data/part*",inferSchema = True)
```

3) After the above step, we create a new dataframe with renamed columns

```
df1 = df.withColumnRenamed("_c0","booking_id")\
            .withColumnRenamed("_c1","customer_id") \
            .withColumnRenamed("_c2","driver_id") \
            .withColumnRenamed("_c3","customer_app_version")  \
            .withColumnRenamed("_c4","customer_phone_os_version") \
            .withColumnRenamed("_c5","pickup_lat") \
            .withColumnRenamed("_c6","pickup_lon") \
            .withColumnRenamed("_c7","drop_lat") \
            .withColumnRenamed("_c8","drop_lon") \
            .withColumnRenamed("_c9","pickup_timestamp")  \
            .withColumnRenamed("_c10","drop_timestamp")  \
            .withColumnRenamed("_c11","trip_fare") \
            .withColumnRenamed("_c12","tip_amount")  \
            .withColumnRenamed("_c13","currency_code") \
            .withColumnRenamed("_c14","cab_color")  \
            .withColumnRenamed("_c15","cab_registration_no") \
            .withColumnRenamed("_c16","customer_rating_by_driver")  \
            .withColumnRenamed("_c17","rating_by_customer")  \
            .withColumnRenamed("_c18","passenger_count")
```

4) Then we create aggregations based on date and store the transformed data in a csv file

```
df2 =  df1.withColumn("booking_date", date_format('pickup_timestamp', "yyyy-MM-dd"))

df2.show(5)
# Date wise aggregation stored in date.
date = df2.select('booking_date').groupBy('booking_date').count()

date.show(10)
date.count()
date.coalesce(1).write.format('com.databricks.spark.csv').mode('overwrite').save('/home/hadoop/datewise_aggreation', header = 'false')
```

The steps to run the python file and the steps are provided in commented file “**step3.pdf**”

Code execution output:
After running the code, we get the following output

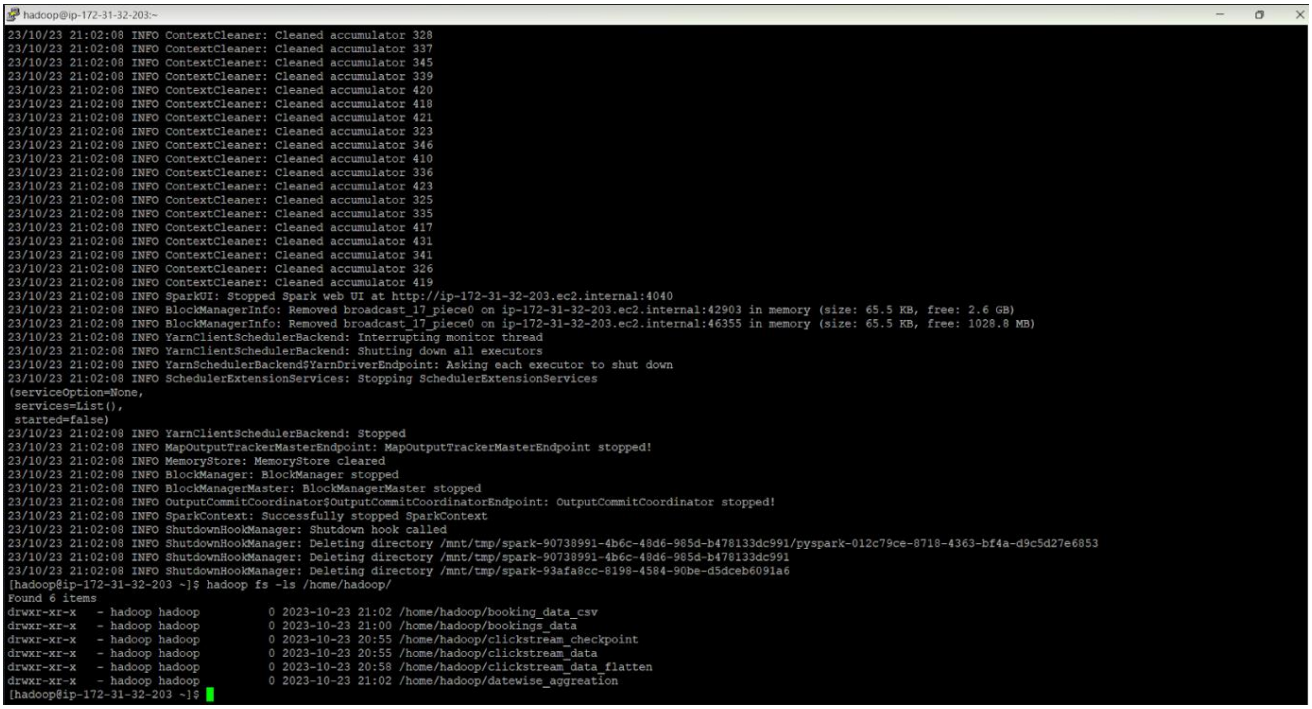

Task 4: Hive managed Table Creation

Please find commented file "**step4.pdf**"

1) Create a Hive-managed table for clickstream data
2) Create a Hive-managed table for bookings data
3) Create a Hive-managed table for aggregated data in Task 3

```
hadoop@ip-172-31-33-247:~
[hadoop@ip-172-31-33-247 ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> create database if not exists cab_booking ;
OK
Time taken: 0.91 seconds
hive> use cab_booking ;
OK
Time taken: 0.073 seconds
hive> create table if not exists clickstream_data
    > (customer_id string ,app_version string, os_version string,lat string ,lon string ,page_id
    > string,button_id string , is_button_click varchar(3) ,is_page_view varchar(3) ,is_scroll_up
    > varchar(3) ,is_scroll_down varchar(3), click_timestamp string )
    > row format delimited fields terminated by ","
    > location '/home/hadoop/clickstream_data_flatten/';
OK
Time taken: 0.422 seconds
hive> create table if not exists booking_data
    > (booking_id string ,customer_id string ,driver_id string , customer_app_version string,
    > customer_phone_os_version string , pickup_lat double , pickup_lon double, drop_lat double,
    > drop_lon double, pickup_timestamp string , drop_timestamp string ,trip_fare int,
    > tip_amount int, currency_code string ,cab_color string, cab_registration_no string ,
    > customer_rating_by_driver int, rating_by_customer int ,passenger_count int )
    > row format delimited fields terminated by ","
    > location '/home/hadoop/booking_data_csv/';
OK
Time taken: 0.07 seconds
hive> create table if not exists datewise_data
    > (booking_date string , count int)
    > row format delimited fields terminated by ","
    > location '/home/hadoop/datewise_aggreation/';
OK
Time taken: 0.05 seconds
hive>
```

Code execution output:

After creating the Hive tables, we check if the data is properly loaded into them

```
hive> use cab_booking;
OK
Time taken: 0.078 seconds
hive> select * from clickstream_data limit 5;
OK
26564820        3.2.35  Android 16.4454865       99.902065       de545711-3914-4450-8c11-b17b8dabb5e1    fcba68aa-1231-11eb-adc1-0242ac120002    No      Yes     No      Yes     "2020-09-14
9:59:07"
31906387        2.4.7   iOS     -64.813749      -133.527040     de545711-3914-4450-8c11-b17b8dabb5e1    a95dd57b-779f-49db-819d-b6960483e554    No      No      Yes     Yes     "2020-05-16
6:30:21"
25713677        3.4.12  Android 89.943435        127.313415      b328829e-17ae-11eb-adc1-0242ac120002    fcba68aa-1231-11eb-adc1-0242ac120002    No      No      Yes     No      "2020-02-09
0:52:13"
83474293        3.1.8   Android -69.939070       -36.451670      e7bc5fb2-1231-11eb-adc1-0242ac120002    e1e99492-17ae-11eb-adc1-0242ac120002    Yes     No      Yes     No      "2020-06-17
0:42:50"
63727807        2.2.9   iOS     64.082108       -81.822078      e7bc5fb2-1231-11eb-adc1-0242ac120002    fcba68aa-1231-11eb-adc1-0242ac120002    No      Yes     Yes     Yes     "2020-07-06
2:51:53"
Time taken: 2.124 seconds, Fetched: 5 row(s)
hive> select * from booking_data limit 5;
OK
BK8968087150    51811359        15055660        2.2.14  Android -49.4319655      103.917851      -58.8043875     146.477367      2020-06-23T19:33:10.000Z        2020-06-06T09:02:10.000Z
34      83      INR     black   054-38-4479     4       3       3
BK629851904     31663218        60872180        3.4.1   iOS     -83.5408405      175.80085       86.20705        128.367238      2020-05-23T12:22:04.000Z        2020-08-09T19:02:56.000Z
26      67      INR     lime    796-39-6801     3       2       4
BK1797410350    86869399        94276051        4.1.36  iOS     -67.8930645      55.234128       -51.1079        -31.07475       2020-05-19T14:14:32.000Z        2020-08-23T18:38:39.000Z
97      63      INR     olive   748-73-1579     1       3       3
BK5788246325    58230837        45457227        2.4.27  Android 13.707887        113.499943      54.3812915      -18.437751      2020-03-24T01:30:15.000Z        2020-05-19T11:16:45.000Z
92      32      INR     white   558-80-6346     3       2       2
BK8342703255    84232510        86494681        4.1.34  Android -6.091461        -114.649789     22.8449505      70.137827       2020-08-03T19:10:52.000Z        2020-03-24T08:25:40.000Z
60      7       INR     blue    068-72-1637     3       3       3
Time taken: 0.16 seconds, Fetched: 5 row(s)
hive> select * from datewise_data limit 5;
OK
2020-06-20      1
2020-04-20      3
2020-05-14      3
2020-04-22      2
2020-03-16      2
Time taken: 0.138 seconds, Fetched: 5 row(s)
hive>
```